

Grandes lignes du projet

Prérequis: - Avoir un compte AWS

- petites notions sur linux, nous travaillerons sur Ubuntu 24.04
- connaissance de base docker, docker-compose

lien, si vous n'avez pas de compte AWS et vous voulez en créer : <https://smart.dhgat.com/step-by-step-guide-to-creating-and-securing-your-aws-account-for-optimal-use/>

un pipeline Airflow orchestrant des extractions de données depuis un bucket S3 de Amazon (weather(météo), COVID) et ensuite des transformations utilisant duckdb via dbt.

- Installation de airflow via le script docker compose que vous pouvez trouver sur le site officiel de airflow: <https://airflow.apache.org/docs/apache-airflow/stable/howto/docker-compose/index.html#fetching-docker-compose-yaml>,

The screenshot shows the Apache Airflow documentation homepage. The top navigation bar includes links for Community, Meetups, Documentation, Use Cases, Announcements, Blog, and Ecosystem. The main content area is titled 'Fetching docker-compose.yaml' with a link. It contains instructions to deploy Airflow on Docker Compose by fetching the docker-compose.yaml file. A command-line example is provided: 'curl -Lfo 'https://airflow.apache.org/docs/apache-airflow/3.1.3/docker-compose.yaml''. A note states that from July 2023 Compose V1 stopped receiving updates, so it's advised to upgrade to a newer version. Below this, a note says the file contains service definitions for various Airflow components like scheduler, dag-processor, api-server, worker, and triggerer. On the right side, there's a sidebar with links for running Airflow in Docker, initializing the environment, and running the CLI commands. A 'Suggest a change on this page' button is also visible.

- cliquez sur le lien en bleu dans l'image ou directement ici : [docker-compose.yml](#) pour télécharger le script docker compose.

- Crée un dossier airflow et mettre le script [docker-compose.yml](#) téléchargé, dans ce dossier

- crée cette structure de base à l'intérieur du dossier airflow avec cette commande : mkdir -p dags
plugins scripts config data/{raw,processed} tests

ensuite :

- on se déplace dans le dossier où se trouve le script et on exécute la commande : docker-compose up -d il installera tout les conteneurs liés à docker et va les lancer.

```
d-a-s@d-a-s-ThinkPad-X380-Yoga:~$ cd Desktop/
d-a-s@d-a-s-ThinkPad-X380-Yoga:~/Desktop$ cd airflow/
d-a-s@d-a-s-ThinkPad-X380-Yoga:~/Desktop/airflow$ docker-compose up
```

The terminal window displays log messages from multiple Airflow DAGs. The logs are color-coded by source: blue for airflow-dag-processor_1, green for airflow-dag-triggerer_1, and red for airflow-apiserver_1. The messages include INFO-level log entries for tasks like 'Sync 1 DAGs', 'Setting next_dagrun for [DAG]', and 'INFO - Sync 2 DAGs'. Specific DAG names mentioned in the logs include 'example_passing_params_via_test_command', 'example_simplest_dag', 'example_branch_datetime_operator', 'example_branch_datetime_operator_2', 'example_branch_datetime_operator_3', 'example_dynamic_task_mapping', and 'tutorial'.

```
d-a-s@d-a-s-ThinkPad-X380-Yoga: ~          d-a-s@d-a-s-ThinkPad-X380-Yoga: ~/Desktop/airflow
example_passing_params_via_test_command to 2025-09-24 21:45:00+00:00, run_after=2025-09-24 21:45:00+00:00
00
airflow-dag-processor_1 | [2025-09-24T21:45:43.387+0000] {dag.py:1622} INFO - Sync 1 DAGs
airflow-dag-processor_1 | [2025-09-24T21:45:43.395+0000] {dag.py:2236} INFO - Setting next_dagrun for example_simplest_dag to None, run_after=None
airflow-dag-processor_1 | [2025-09-24T21:45:43.461+0000] {dag.py:1622} INFO - Sync 3 DAGs
airflow-dag-processor_1 | [2025-09-24T21:45:43.474+0000] {dag.py:2236} INFO - Setting next_dagrun for example_branch_datetime_operator to 2025-09-24 00:00:00+00:00, run_after=2025-09-24 00:00:00+00:00
airflow-dag-processor_1 | [2025-09-24T21:45:43.479+0000] {dag.py:2236} INFO - Setting next_dagrun for example_branch_datetime_operator_2 to 2025-09-24 00:00:00+00:00, run_after=2025-09-24 00:00:00+00:00
airflow-dag-processor_1 | [2025-09-24T21:45:43.483+0000] {dag.py:2236} INFO - Setting next_dagrun for example_branch_datetime_operator_3 to 2025-09-24 00:00:00+00:00, run_after=2025-09-24 00:00:00+00:00
airflow-dag-processor_1 | [2025-09-24T21:45:43.536+0000] {dag.py:1622} INFO - Sync 2 DAGs
airflow-dag-processor_1 | [2025-09-24T21:45:43.542+0000] {dag.py:2236} INFO - Setting next_dagrun for example_dynamic_task_mapping to None, run_after=None
airflow-dag-processor_1 | [2025-09-24T21:45:43.544+0000] {dag.py:2236} INFO - Setting next_dagrun for example_task_mapping_second_order to None, run_after=None
airflow-dag-processor_1 | [2025-09-24T21:45:43.595+0000] {dag.py:1622} INFO - Sync 1 DAGs
airflow-dag-processor_1 | [2025-09-24T21:45:43.615+0000] {dag.py:2236} INFO - Setting next_dagrun for example_branch_dop_operator_v3 to 2025-09-24 21:45:00+00:00, run_after=2025-09-24 21:45:00+00:00
airflow-dag-processor_1 | [2025-09-24T21:45:43.672+0000] {dag.py:1622} INFO - Sync 1 DAGs
airflow-dag-processor_1 | [2025-09-24T21:45:43.682+0000] {dag.py:2236} INFO - Setting next_dagrun for tutorial to 2025-09-24 21:45:43.681904+00:00, run_after=2025-09-24 21:45:43.681904+00:00
```

The terminal window shows the output of the 'docker ps' command, listing all running Docker containers. The table includes columns for CONTAINER ID, IMAGE, COMMAND, CREATED, and STATUS. The STATUS column indicates the health of each container, with most being 'Up' and one being 'healthy'. The containers listed are airflow_airflow-worker_1, airflow_airflow-triggerer_1, airflow_airflow-dag-processor_1, airflow_apiserver_1, airflow_airflow-scheduler_1, airflow_postgres_1, and airflow_redis_1.

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
47e35a0c0722	airflow_airflow-worker	" /usr/bin/dumb-init ... "	24 hours ago	Up About a minut
e (healthy)	8080/tcp		airflow_airflow-worker_1	
6c04087dd2e0	airflow_airflow-triggerer	" /usr/bin/dumb-init ... "	24 hours ago	Up About a minut
e (healthy)	8080/tcp		airflow_airflow-triggerer_1	
d79c5bdb9c5d	airflow_airflow-dag-processor	" /usr/bin/dumb-init ... "	24 hours ago	Up About a minut
e (healthy)	8080/tcp		airflow_airflow-dag-processor_1	
409bfccb5407	airflow_airflow-apiserver	" /usr/bin/dumb-init ... "	24 hours ago	Up About a minut
e (healthy)	0.0.0.0:8080->8080/tcp, [::]:8080->8080/tcp		airflow_airflow-apiserver_1	
88bc4d21f31e	airflow_airflow-scheduler	" /usr/bin/dumb-init ... "	24 hours ago	Up About a minut
e (healthy)	8080/tcp		airflow_airflow-scheduler_1	
cde6b682c55d	postgres:13	" docker-entrypoint.s... "	24 hours ago	Up About a minut
e (healthy)	5432/tcp		airflow_postgres_1	
e3754f863416	redis:7.2-bookworm	" docker-entrypoint.s... "	24 hours ago	Up About a minut
e (healthy)	6379/tcp		airflow_redis_1	

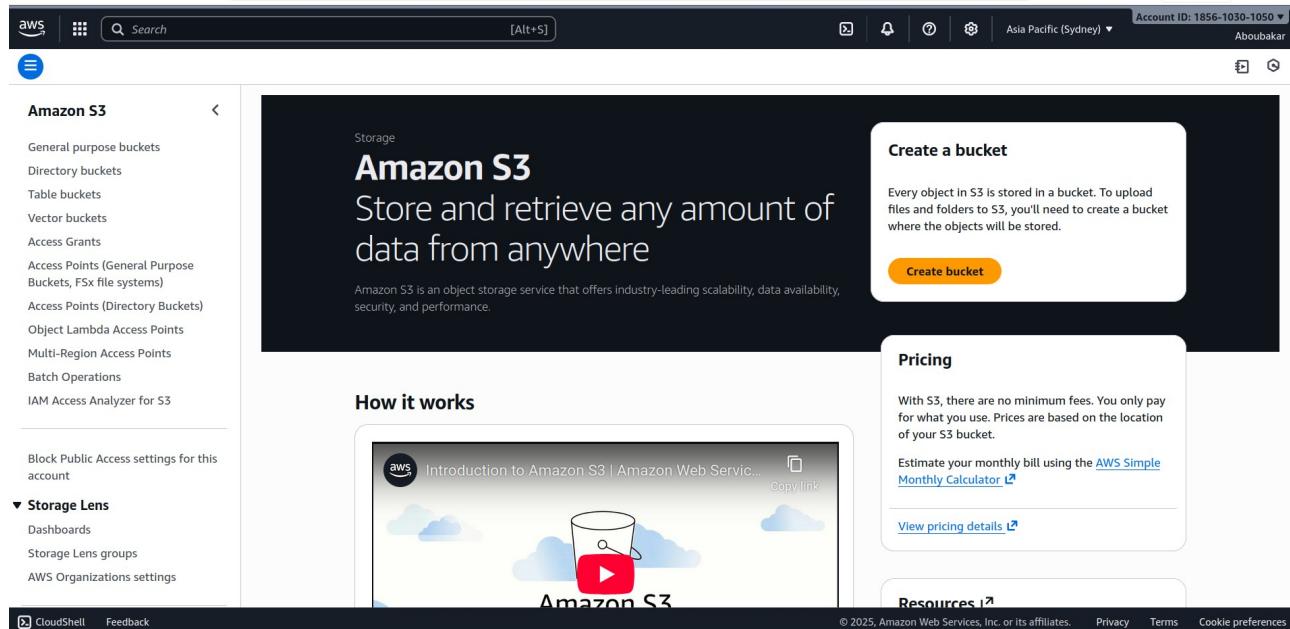
- Installer dbt-core + l'adaptateur DuckDB car on ne va pas utiliser Postgres mais DuckDB vu que

S3 est un outil de stockage objet qui permet de faire des requêtes de données qui n'ont pas une structure relationnelle pour pouvoir effectuer des requêtes sur S3, on ne peut pas utiliser Postgres car gérant les db relationnelles, on va utiliser DuckDB qui peut requêter sur des fichiers Json, CSV et à des connecteurs pour S3 qu'on va utiliser.

On le fera plus tard vers la fin dans un conteneur airflow pour qu'il soit visible à l'intérieur d'airflow.

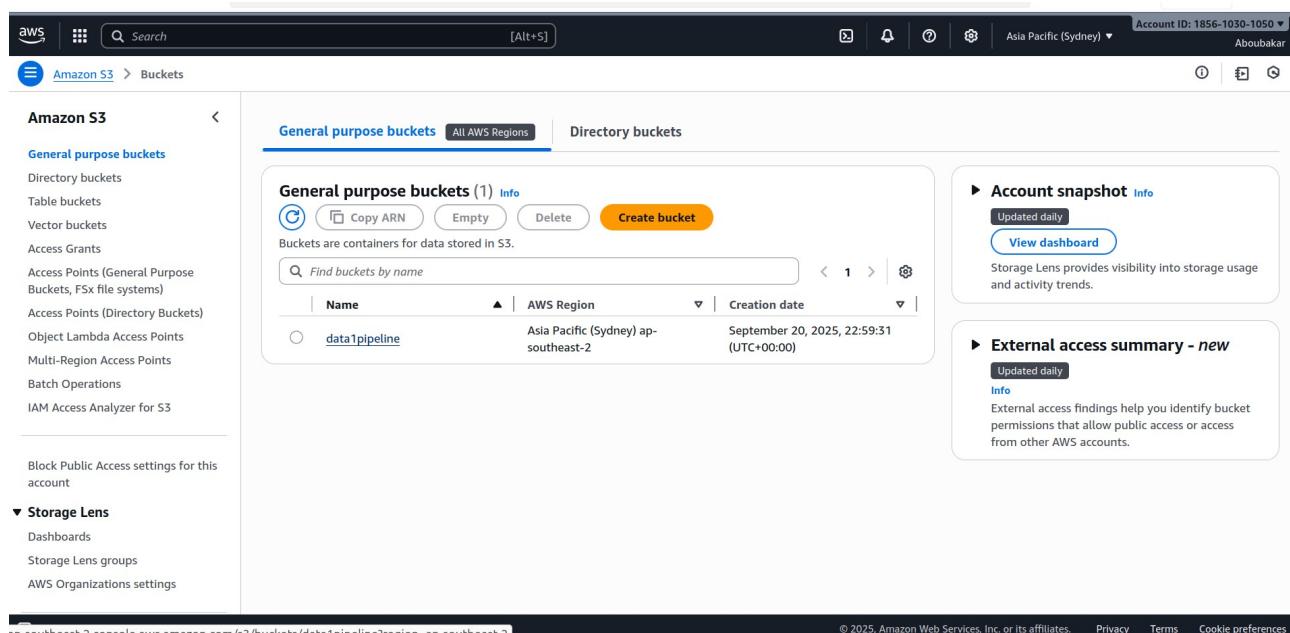
- Créez un bucket S3, on l'a nommé dans notre cas data1pipeline vous pouvez le changer

- Allez sur votre console AWS, connectez-vous, tapez dans la barre de recherche S3, vous allez être ici



The screenshot shows the AWS S3 landing page. On the left, there's a sidebar with navigation links like 'General purpose buckets', 'Storage Lens', and 'CloudShell'. The main content area features a large 'Amazon S3' heading with the subtext 'Store and retrieve any amount of data from anywhere'. Below this is a brief description of Amazon S3's benefits. To the right, there are three callout boxes: 'Create a bucket', 'Pricing', and 'Resources'. The 'Create a bucket' box contains a button labeled 'Create bucket'. The 'Pricing' box states that there are no minimum fees and provides a link to the 'AWS Simple Monthly Calculator'. The 'Resources' box links to various AWS services. At the bottom of the page, there's a footer with copyright information and links to 'Privacy', 'Terms', and 'Cookie preferences'.

- Cliquez sur general purpose bucket



The screenshot shows the 'General purpose buckets' page. The sidebar on the left is identical to the previous screenshot. The main content area displays a table for 'General purpose buckets'. It shows one bucket named 'data1pipeline' with the following details: Name (data1pipeline), AWS Region (Asia Pacific (Sydney) ap-southeast-2), and Creation date (September 20, 2025, 22:59:31 (UTC+00:00)). There are buttons for 'Copy ARN', 'Empty', 'Delete', and 'Create bucket'. To the right of the table, there are two callout boxes: 'Account snapshot' and 'External access summary - new'. Both boxes have 'Updated daily' status indicators and provide links to their respective dashboards. The footer at the bottom includes the URL 'ap-southeast-2.console.aws.amazon.com/s3/buckets/data1pipeline?region=ap-southeast-2', along with 'Privacy', 'Terms', and 'Cookie preferences' links.

- Cliquez sur create bucket et entrez le nom laissez les parametre ar defaut et cliquez tout en bas sur create bucket

The screenshot shows the AWS S3 'Create bucket' interface. In the 'General configuration' section, the 'Bucket name' is set to 'amzn-s3-demo-bucket'. Under 'Object Ownership', 'ACLs disabled (recommended)' is selected. The status bar at the bottom indicates 'CloudShell' and 'Feedback'.

The screenshot shows the 'Default encryption' section where 'Server-side encryption with Amazon S3 managed keys (SSE-S3)' is selected. Below it, under 'Bucket Key', 'Enable' is selected. A note at the bottom says 'After creating the bucket, you can upload files and folders to the bucket, and configure additional bucket settings.' The status bar at the bottom indicates 'CloudShell' and 'Feedback'.

The screenshot shows the AWS S3 buckets page. At the top, there's a green banner with the message "Successfully created bucket 'data2pipeline'". Below the banner, there are tabs for "General purpose buckets" (selected) and "All AWS Regions". A search bar and a "Create bucket" button are also present. On the left, a table lists two buckets: "data1pipeline" and "data2pipeline", both in the "Asia Pacific (Sydney) ap-southeast-2" region. To the right, there are two cards: "Account snapshot" (updated daily) and "External access summary - new" (updated daily). The bottom of the page includes links for CloudShell, Feedback, Privacy, Terms, and Cookie preferences.

- Crée S3 Access Key et Secret Key

- 1- Ces identifiants permettent à dbt (ou tout autre outil) d'accéder à votre bucket S3.
- 2- Ils sont générés depuis la console AWS IAM (Identity and Access Management) :
- 3- Créer un utilisateur IAM avec accès programmatique (Programmatic access).
- 4- Attacher une politique (Policy) permettant l'accès à S3 (ex: AmazonS3FullAccess pour test ou une policy restreinte pour la production).
- 6- Cliquez sur votre utilisateur et cherchez : create access key et générer les
- 7- La console génère alors une Access Key ID et une Secret Access Key.
Ces clés peuvent être utilisées dans dbt, Airflow, ou tout script pour se connecter à S3.

- Placer les fichiers data_extraction_dag.py et dbt_pipeline_dag.py dans le dossier airflow créé par docker compose

- Crée dans le dossier airflow, le dossier dbt et à l'intérieur de ce dossier, le dossier "votre projet"

```
# Aller dans ton dossier airflow (si tu es sur ton Desktop par exemple)
cd ~/Desktop/airflow
```

```
# Créer le dossier dbt à l'intérieur du dossier airflow
mkdir -p dbt
```

```
# Aller dans ce dossier dbt
cd dbt
```

```
# Créer le dossier pour ton projet dbt (exemple : my_dbt_project)
mkdir -p my_dbt_project
```

```
# Initialiser un projet dbt à l'intérieur
cd my_dbt_project
dbt init my_dbt_project
```

Volumes montés dans le conteneur Airflow via Docker Compose :

- /home/d-a-s/Desktop/airflow/dbt:/opt/airflow/dbt: on le met dans la partie volume de notre fichier docker compose

- Le dossier local ~/airflow/dbt est monté dans le conteneur à /opt/airflow/dbt.
- Sert à stocker les projets dbt (models, macros, analyses, seeds, etc.) de manière persistante.
- Toute modification dans le conteneur sera visible sur la machine hôte et vice-versa.

- /home/d-a-s/.dbt:/opt/airflow/.dbt : on le met dans la partie volume de notre fichier docker compose
 - Le dossier local ~/.dbt, il contient le fichier profile, vous pouvez le mettre ailleurs mais il faut spécifier le dossier et monté dans le conteneur à /opt/airflow/.dbt.
 - Contient les fichiers de configuration dbt, notamment profiles.yml qui définit les connexions aux bases de données.
 - Permet au conteneur d'utiliser les mêmes configurations que celles définies sur la machine hôte.

Configuration dbt pour le projet my_dbt_project (profiles.yml) :

```

- target: dev
  → Définit l'environnement de travail par défaut pour dbt.
- outputs:
  dev:
    type: duckdb
    path: /opt/airflow/dbt/my_dbt_project/my_duckdb.duckdb
    → Fichier DuckDB stocké dans le dossier du projet dbt monté dans le conteneur Airflow.
    → Persistant : toutes les modifications restent sur la machine hôte via le volume Docker.
  extensions:
    - httpfs
    → Permet à DuckDB d'accéder à des fichiers stockés sur S3 directement depuis dbt.
  threads: 1
  → Nombre de threads utilisés pour l'exécution des modèles dbt.
  schema: main
  → Schéma par défaut utilisé dans DuckDB.
  s3_region: ap-southeast-2
  s3_access_key_id: AKIASWNZ4NZ5JMMRML0V
  s3_secret_access_key: KTFHVIM04bmzf7Dw85KfeYDd8jO+5Hi4PiV59cDi
  → Informations de connexion pour accéder au bucket S3 depuis dbt/DuckDB.

```

Fichier sources.yml (dans models/staging)

Rôle :

- Définit les sources de données externes que dbt va utiliser dans les modèles.
- Permet à dbt de référencer les fichiers S3 (ou d'autres sources) de manière centralisée et documentée.
- Facilite la traçabilité et la gestion des données avant toute transformation (staging).

Contenu et explication :

version: 2

sources:

- name: raw

```

# Nom logique de la source de données
tables:
- name: weather_data # Table ou fichier météo
description: "Fichiers météo depuis S3"
external:
# Données externes (non transformées)
location: "s3://data1pipeline/raw/weather/" # Chemin S3
format: json
# Format des fichiers
pattern: "*.json" # Pattern pour filtrer les fichiers
- name: covid_data # Table ou fichier COVID
description: "Fichiers COVID depuis S3"
external:
location: "s3://data1pipeline/raw/covid/"
format: csv
pattern: "*.csv"

```

Résumé :

- Ce fichier n'est pas strictement obligatoire pour faire fonctionner dbt, mais il est fortement recommandé.
 - Il sert à documenter les sources et permettre aux modèles dbt de les référencer facilement via `source('raw', 'weather_data')` ou `source('raw', 'covid_data')`.
 - Cela centralise la configuration S3, donc si le chemin change, il suffit de modifier ici.
- Étapes pour ajouter tes modèles

1. Aller dans ton projet dbt :

```
cd ~/Desktop/airflow/dbt/my_dbt_project
```

2. Créer les dossiers nécessaires :

```
mkdir -p models/staging
```

```
mkdir -p models/marts
```

mettez y les fichiers

```
-- models/staging/stg_weather.sql
```

```
-- models/staging/stg_covid.sql
```

```
-- models/mart/daily_weather_metrics.sqlCréation et récupération de la clé API OpenWeatherMap
```

1. Créer un compte :

Aller sur https://home.openweathermap.org/users/sign_up

2. Connexion :

Se connecter à <https://home.openweathermap.org>

3. Générer une clé API :

- Aller dans le menu API keys.
- Cliquer sur "Generate" ou utiliser la clé par défaut fournie.
- Donner un nom à la clé (exemple : airflow_project_key).

Vous utiliserez cette clé dans le data_extraction_dag.py pour récupérer les données de meteo(weather)

4.

1 Accéder au conteneur Airflow

Identifie d'abord le nom du conteneur Airflow en cours d'exécution :

avec docker ps

Ensuite, connecte-toi au conteneur (par ex. le scheduler ou le webserver)

avec la commande : docker exec -it airflow-scheduler-1 bash

5.

2 Installer dbt et dbt avec l adaptateur duckdb dans le conteneur avec : pip install --no-cache-dir dbt-core dbt-duckdb

6. Vérifier l'installation avec

dbt –version

vous dois voir quelque chose comme :

installed version: 1.8.x

plugins:

- duckdb: 1.8.x

7. ouvrir la page web de airflow http://localhost:8080/

ET executé les dag : data_extraction_dag et dbt_pipeline_dag

Resultats

The screenshot shows the AWS S3 console interface. At the top, the URL is `ap-southeast-2.console.aws.amazon.com/s3/buckets/data1pipeline?region=ap-southeast-2&bucketType=general&prefix=raw%2F&tab=Objects`. The account ID is 1856-1030-1050, and the region is Asia Pacific (Sydney). The 'raw/' folder is selected. The 'Objects' tab is active, showing one object: 'api_data/' (Folder). There are buttons for Actions, Create folder, and Upload.

CloudShell Feedback © 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

On vérifie qu'il y a des données dans le bucket

The screenshot shows the Airflow UI at `http://localhost:8080/dags/data_extraction_pipeline`. The left sidebar has tabs for Home, Dags, Assets, Browse, Admin, Security, Docs, and User. The 'Dags' tab is selected, showing the 'data_extraction_pipeline' DAG. The DAG visualization shows tasks: start, extract_covid..., extract_weather..., and end. A red bar indicates a failed task. The DAG details show it was scheduled at 1:00:00 and the latest run was 2025-09-24, 21:44:29. The next run is scheduled for 2025-09-24, 22:44:29. The owner is airflow. The 'Overview' tab is selected, showing a timeline from 2025-09-23, 21:56:23 to 2025-09-24, 21:56:23. It highlights two failed tasks and a failed run. Below this is a chart titled 'Last 12 Dag Runs' showing duration in seconds. The 'Recent Failed Task Logs' section shows an error log for the 'extract_covid_data' task on 2025-09-24, 21:44:29, with the error message: 'Task failed with exception: source="task" URLError: <urlopen error [Errno -3] Temporary failure in name resolution>'.

Le pipeline d'extraction des données des API et le stockage vers S3

Trigger DAG - data_extraction_pipeline

Single Run
Trigger a single run of this DAG

Backfill
Run this DAG for a range of dates

Last 12 Dag Runs

Recent Failed Task Logs

extract_covid_data 2025-09-24, 21:44:29

8 [2025-09-24, 21:50:16] ERROR - Task failed with exception: source="task"
▼ URLLError: <urlopen error [Errno -3] Temporary failure in name resolution>

On lance le pipeline d'extraction de données des API de Covid et Météo et on envoie dans le bucket S3 Selon le script data_extraction_dag.py qui dans dags du dossier airflow

Dag Run 2025-09-24, 21:56:32 ✓ success

Logical Date	Run Type	Start	End	Duration	Dag Version(s)
2025-09-24, 21:56:32	▶ manual	2025-09-24, 21:56:39	2025-09-24, 21:57:27	00:00:48	v35

Task Instances

Task ID	Map Index	State	Start Date	End Date	Try Number	Operator	Duration	Dag V
end		✓ success	2025-09-24, 21:57:26	2025-09-24, 21:57:26	1	EmptyOperator	00:00:00	v35
extract_covid_data		✓ success	2025-09-24, 21:56:40	2025-09-24, 21:57:26	1	PythonOperator	00:00:45	v35
extract_weather_data		✓ success	2025-09-24, 21:56:40	2025-09-24, 21:56:49	1	PythonOperator	8.63s	v35
start		✓ success	2025-09-24, 21:56:39	2025-09-24, 21:56:39	1	EmptyOperator	00:00:00	v35

Le pipeline exécuté avec succès

Sep 24 21:57

aws | [Alt+S] | Asia Pacific (Sydney) | Account ID: 1856-1030-1050 | Aboubakar

Amazon S3 > Buckets > data1pipeline > raw/

raw/

Objects Properties

Objects (3)

Objects are fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

Name	Type	Last modified	Size	Storage class
api_data/	Folder	-	-	-
covid/	Folder	-	-	-
weather/	Folder	-	-	-

Copy S3 URI

CloudShell Feedback © 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Sous bucket créés

Sep 24 21:57

aws | [Alt+S] | Asia Pacific (Sydney) | Account ID: 1856-1030-1050 | Aboubakar

Amazon S3 > Buckets > data1pipeline > raw/ > weather/

weather/

Objects (1)

Objects are fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

Name	Type	Last modified	Size	Storage class
Paris_20250924_215643.json	json	September 24, 2025, 21:56:48 (UTC+00:00)	566.0 B	Standard

Copy S3 URI

CloudShell Feedback © 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

données dans le sous bucket weather

The screenshot shows the AWS S3 console interface. At the top, the URL is ap-southeast-2.console.aws.amazon.com/s3/buckets/data1pipeline?region=ap-southeast-2&bucketType=general&prefix=raw/covid/&short. The account ID is 1856-1030-1050, and the region is Asia Pacific (Sydney). The bucket name is 'data1pipeline'. The path is 'raw/covid/'. A single object named 'covid/' is listed, which contains the file '20250924_215704.csv'. The file is a CSV type, 8.6 MB in size, and has a standard storage class. It was last modified on September 24, 2025, at 21:57:11 (UTC+00:00).

The screenshot shows the dbt Data Pipeline interface. The left sidebar has navigation links: Home, Dags, Assets, Browse, Admin, Security, Docs, and User. The 'Dags' link is selected. In the center, a specific DAG named 'dbt_data_pipeline' is shown with a run history for '2025-09-24, 00:00:00'. The run status is 'success'. The table below lists the task instances for this run:

Task ID	Map Index	State	Start Date	End Date	Try Number	Operator	Duration	Dag Version
end		✓ success	2025-09-24, 22:00:00	2025-09-24, 22:00:00	1	EmptyOperator	00:00:00	v1
send_success_notification		✓ success	2025-09-24, 21:59:59	2025-09-24, 21:59:59	1	PythonOperator	0.65s	v1
data_quality_check		✓ success	2025-09-24, 21:59:57	2025-09-24, 21:59:58	1	PythonOperator	0.77s	v1
dbt_test		✓ success	2025-09-24, 21:59:47	2025-09-24, 21:59:56	1	BashOperator	8.88s	v1
dbt_run		✓ success	2025-09-24, 21:58:34	2025-09-24, 21:59:46	3	BashOperator	00:01:12	v1
start		✓ success	2025-09-24, 21:58:33	2025-09-24, 21:58:33	2	EmptyOperator	00:00:00	v1

Execution du pipeline de transformation dbt des données stocker dans le sous bucket covid et weather sur S3

The screenshot shows the Airflow web interface. On the left is a sidebar with icons for Home, Dags, Assets, Browse, Admin, Security, Docs, and User. The main area displays a dag named 'dbt_data_pipeline' with a single task named 'dbt_run'. The task status is 'success' at '2025-09-24, 21:58:34'. The logs tab shows the command-line output of the dbt run, which includes multiple INFO log entries indicating successful database operations. The logs end with a message about pushing XCom data.

```

16 [2025-09-24, 21:58:44] INFO - [0m21:58:44: source="airflow.task-hooks.airflow.providers.standard_hooks.subprocess.SubprocessHook"]
17 [2025-09-24, 21:58:44] INFO - [0m21:58:44 1 of 5 START sql table model main.my_first_dbt_model ..... [32mOK[0m in 0.33s]: soi
18 [2025-09-24, 21:58:45] INFO - [0m21:58:45 1 of 5 OK created sql table model main.my_first_dbt_model ..... [32mOK[0m in 0.33s]: soi
19 [2025-09-24, 21:58:45] INFO - [0m21:58:45 2 of 5 START sql table model main.stg_covid ..... [32mOK[0m in 0.33s]: soi
20 [2025-09-24, 21:59:28] INFO - [0m21:59:28 2 of 5 OK created sql table model main.stg_covid ..... [32mOK[0m in 43.19s]: soi
21 [2025-09-24, 21:59:28] INFO - [0m21:59:28 3 of 5 START sql view model main_staging.stg_weather ..... [32mOK[0m in 43.19s]: soi
22 [2025-09-24, 21:59:33] INFO - [0m21:59:33 3 of 5 OK created sql view model main_staging.stg_weather ..... [32mOK[0m in 5.22s]: soi
23 [2025-09-24, 21:59:33] INFO - [0m21:59:33 4 of 5 START sql view model main.my_second_dbt_model ..... [32mOK[0m in 0.10s]: soi
24 [2025-09-24, 21:59:33] INFO - [0m21:59:33 4 of 5 OK created sql view model main.my_second_dbt_model ..... [32mOK[0m in 0.10s]: soi
25 [2025-09-24, 21:59:33] INFO - [0m21:59:33 5 of 5 START sql table model main.mart.daily_weather_metrics ..... [32mOK[0m in 10.76s]: soi
26 [2025-09-24, 21:59:44] INFO - [0m21:59:44 5 of 5 OK created sql table model main.mart.daily_weather_metrics ..... [32mOK[0m in 10.76s]: soi
27 [2025-09-24, 21:59:44] INFO - [0m21:59:44 Finished running 3 table models, 2 view models in 0 hours 1 minutes and 0.49 seconds (60.49s).: source="airflow.task-hooks.airflow.providers.standard_hooks.subprocess.SubprocessHook"
28 [2025-09-24, 21:59:44] INFO - [0m21:59:44 [32mCompleted successfully[0m: source="airflow.task-hooks.airflow.providers.standard_hooks.subprocess.SubprocessHook"]
29 [2025-09-24, 21:59:44] INFO - [0m21:59:44 [source="airflow.task-hooks.airflow.providers.standard_hooks.subprocess.SubprocessHook"]
30 [2025-09-24, 21:59:44] INFO - [0m21:59:44 [32mCompleted successfully[0m: source="airflow.task-hooks.airflow.providers.standard_hooks.subprocess.SubprocessHook"]
31 [2025-09-24, 21:59:44] INFO - [0m21:59:44 [source="airflow.task-hooks.airflow.providers.standard_hooks.subprocess.SubprocessHook"]
32 [2025-09-24, 21:59:44] INFO - [0m21:59:44 Done. PASS=0 WARN=0 ERROR=0 SKIP=0 NO_OP=0 TOTAL=5: source="airflow.task-hooks.airflow.providers.standard_hooks.subprocess.SubprocessHook"
33 [2025-09-24, 21:59:46] INFO - [0m21:59:46 Command exited with return code 0: source="airflow.task-hooks.airflow.providers.standard_hooks.subprocess.SubprocessHook"
34 [2025-09-24, 21:59:46] INFO - [0m21:59:46 Pushing xcom: ti="RuntimetimeTaskInstance(id='01997dbc-3f70-7f85-acf2-95984ac341c4', task_id='dbt_run', dag_id='dbt_dag')": source="airflow.task-hooks.airflow.providers.standard_hooks.subprocess.SubprocessHook"

```

Tout les modèles sont exécutés

The terminal session shows the user executing dbt commands within an airflow scheduler container. The user runs 'docker exec -it airflow_airflow-scheduler_1 bash' and then enters the container. Inside, they run 'cd /opt/airflow/dbt/my_dbt_project' and 'bt debug --profiles-dir /opt/airflow/.dbt --profile my_dbt_project'. The output shows the dbt debug process starting and listing profiles, adapters, configurations, and connection details for the 'my_duckdb' database.

```

d-a-s@d-a-s-ThinkPad-X380-Yoga:~/Desktop/airflow
d-a-s@d-a-s-ThinkPad-X380-Yoga:~/Desktop/airflow
d-a-s@d-a-s-ThinkPad-X380-Yoga:~/Desktop/airflow$ docker exec -it airflow_airflow-scheduler_1 bash
root@88bc4d21f31e:/opt/airflow# cd /opt/airflow/dbt/my_dbt_project
root@88bc4d21f31e:/opt/airflow/dbt/my_dbt_project# bt debug --profiles-dir /opt/airflow/.dbt --profile my_dbt_project
bash: bt: command not found
root@88bc4d21f31e:/opt/airflow/dbt/my_dbt_project# dbt debug --profiles-dir /opt/airflow/.dbt --profile my_dbt_project
22:03:27 Running with dbt=1.10.11
22:03:28 dbt version: 1.10.11
22:03:28 python version: 3.12.11
22:03:28 python path: /home/airflow/.local/bin/python
22:03:28 os info: Linux-6.14.0-24-generic-x86_64-with-glibc2.36
22:03:28 Using profiles dir at /opt/airflow/.dbt
22:03:28 Using profiles.yml file at /opt/airflow/.dbt/profiles.yml
22:03:28 Using dbt_project.yml file at /opt/airflow/dbt/my_dbt_project/dbt_project.yml
22:03:28 adapter type: duckdb
22:03:28 adapter version: 1.9.6
22:03:28 Configuration:
22:03:28   profiles.yml file [OK found and valid]
22:03:28   dbt_project.yml file [OK found and valid]
22:03:28 Required dependencies:
22:03:28   - git [OK found]
22:03:28 Connection:
22:03:28   database: my_duckdb
22:03:28   schema: main
22:03:28   path: /opt/airflow/dbt/my_dbt_project/my_duckdb.duckdb
22:03:28   config_options: None
22:03:28   extensions: ['httpfs']
22:03:28   settings: {}
22:03:28   external_root: .
22:03:28   use_credential_provider: None
22:03:28   attach: None
22:03:28   filesystems: None
22:03:28   remote: None
22:03:28   plugins: None
22:03:28   disable_transactions: False

```

On se connecte au conteneur airflow scheduler où on a installé dbt et on exécute la commande pour voir que la connexion de DuckDB à S3 est fait

```

Sep 24 22:04 68%
d-a-s@d-a-s-ThinkPad-X380-Yoga: ~/Desktop/airflow
d-a-s@d-a-s-ThinkPad-X380-Yoga: ~/Desktop/airflow
d-a-s@d-a-s-ThinkPad-X380-Yoga: ~/Desktop/airflow

22:03:27 Running with dbt=1.10.11
22:03:28 dbt version: 1.10.11
22:03:28 python version: 3.12.11
22:03:28 python path: /home/airflow/.local/bin/python
22:03:28 os info: Linux-6.14.0-24-generic-x86_64-with-glibc2.36
22:03:28 Using profiles dir at /opt/airflow/.dbt
22:03:28 Using profiles.yml file at /opt/airflow/.dbt/profiles.yml
22:03:28 Using dbt_project.yml file at /opt/airflow/dbt/my_dbt_project/dbt_project.yml
22:03:28 adapter type: duckdb
22:03:28 adapter version: 1.9.6
22:03:28 Configuration:
22:03:28   profiles.yml file [OK found and valid]
22:03:28   dbt_project.yml file [OK found and valid]
22:03:28 Required dependencies:
22:03:28   - git [OK found]

22:03:28 Connection:
22:03:28   database: my_duckdb
22:03:28   schema: main
22:03:28   path: /opt/airflow/dbt/my_dbt_project/my_duckdb.duckdb
22:03:28   config_options: None
22:03:28   extensions: ['httpfs']
22:03:28   settings: {}
22:03:28   external_root: .
22:03:28   use_credential_provider: None
22:03:28   attach: None
22:03:28   filesystems: None
22:03:28   remote: None
22:03:28   plugins: None
22:03:28   disable_transactions: False
22:03:28 Registered adapter: duckdb=1.9.6
22:03:54 Connection test: [OK connection ok]

22:03:54 All checks passed!
root@88bc4d21f31e:/opt/airflow/dbt/my_dbt_project# 

```

Connexion réussi

```

root@88bc4d21f31e:/opt/airflow/dbt/my_dbt_project# curl -L https://github.com/duckdb/duckdb/releases/download/v1.1.3/duckdb_cli-linux-amd64.zip -o duckdb.zip
% Total    % Received % Xferd  Average Speed   Time   Time  Current
          Dload  Upload Total Spent   Left Speed
0     0    0     0    0     0      0 --:--:-- --:--:-- --:--:-- 0
100 15.9M 100 15.9M    0     0  1466k      0:00:11  0:00:11  0:00:00 2206k
root@88bc4d21f31e:/opt/airflow/dbt/my_dbt_project# apt-get install -y unzip
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
Unzip is already the newest version (6.0-28).
0 upgraded, 0 newly installed, 0 to remove and 66 not upgraded.
root@88bc4d21f31e:/opt/airflow/dbt/my_dbt_project# unzip duckdb.zip -d /usr/local/bin
chmod +x /usr/local/bin/duckdb
Archive:  duckdb.zip
  inflating: /usr/local/bin/duckdb
root@88bc4d21f31e:/opt/airflow/dbt/my_dbt_project# duckdb --version
v1.1.3 19864453f7
root@88bc4d21f31e:/opt/airflow/dbt/my_dbt_project# duckdb my_duckdb.duckdb
Error: unable to open database "my_duckdb.duckdb": IO Error: Could not set lock on file "my_duckdb.duckdb": Conflicting lock is held in /usr/local/bin/python3.12 (PID 963). See also https://duckdb.org/docs/connect/concurrency
root@88bc4d21f31e:/opt/airflow/dbt/my_dbt_project# ps -p 963 -f

```

Installation de DuckDB dans le conteneur pour vérifier qu'il y a effectivement les tables liés à l'exécution de dbt run

```

bash: ps: command not found
root@88bc4d21f31e:/opt/airflow/dbt/my_dbt_project# kill -9 963
root@88bc4d21f31e:/opt/airflow/dbt/my_dbt_project# duckdb my_duckdb.duckdb -readonly
v1.1.3 19864453f7
Enter ".help" for usage hints.
D .show tables
Usage: .show
D SHOW TABLES;
+-----+
| name        |
| varchar     |
+-----+
| my_first_dbt_model |
| my_second_dbt_model |
| stg_covid    |
+-----+
D -- Voir tous les schémas
D .show schemas;
Usage: .show
D
D -- Voir toutes les tables et vues dans tous les schémas
D SELECT table_schema, table_name, table_type
  FROM information_schema.tables
  WHERE table_schema IN ('main', 'main_staging', 'main_mart');
+-----+-----+-----+
| table_schema | table_name | table_type |
| varchar     | varchar    | varchar    |
+-----+-----+-----+
| main        | my_first_dbt_model | BASE TABLE |
| main        | stg_covid       | BASE TABLE |
| main_mart   | daily_weather_metrics | BASE TABLE |
| main        | my_second_dbt_model | VIEW       |
| main_staging | stg_weather    | VIEW       |
+-----+-----+-----+

```

Verification qu'on a toutes les tables

date	country	province	confirmed	deaths	recovered
2020-01-22	Afghanistan		0	0	0.0
2020-01-23	Afghanistan		0	0	0.0
2020-01-24	Afghanistan		0	0	0.0
2020-01-25	Afghanistan		0	0	0.0
2020-01-26	Afghanistan		0	0	0.0
2020-01-27	Afghanistan		0	0	0.0
2020-01-28	Afghanistan		0	0	0.0
2020-01-29	Afghanistan		0	0	0.0
2020-01-30	Afghanistan		0	0	0.0
2020-01-31	Afghanistan		0	0	0.0

10 rows 6 columns

```
D SELECT * FROM main_stg_covid LIMIT 10;
```

city	date	avg_temperature	max_temperature	min_temperature	avg_humidity	observations_count
"Paris"	2025-09-24	10.55	10.55	10.55	93.0	1

données dans staging covid et metrics de météo

```
D SELECT * FROM main_staging.stg_weather LIMIT 2;
```

100%

city	observation_time	temperature	humidity	weather_description	wind_speed	loaded_at
"Paris"	2025-09-24 21:50:41+00	10.55	93.0	"overcast clouds"	4.63	2025-09-24 22:56:08.078+00
"Paris"	2025-09-24 22:36:56+00	10.55	93.0	"overcast clouds"	3.6	2025-09-24 22:56:08.078+00

données dans stg_weather